

Pressure-Volume Loop Shape as a Fatigue Leading Indicator for Pressure-Only Proprioception in Shared-Manifold Soft Grippers: A Simulation Study

Author: Mergen Ulziibayar · NYU Tandon, Dept. of Mechanical & Aerospace Engineering

Status: Draft v1 (2026-06-24). **This is a simulation-only modeling study.** No physical experiments are reported; every result is synthetic and is labeled as such. Modeling assumptions are stated as choices, not as calibrated predictions for a physical actuator.

Venue note (decision pending, not a content blocker). This draft is written venue-neutral. Because the contribution is simulation-only, the honest first destination is an **arXiv preprint**. RA-L and *Soft Robotics* (the proposal's targets) typically expect hardware validation; they become appropriate only if/when a physical bench campaign is added. Do not submit anywhere without an explicit go decision.

Abstract

Soft pneumatic grippers face two coupled deployment problems: silicone chambers fail by cyclic fatigue, and recovering gripper pose without added sensors relies on the pressure signals already in the control loop. We study, in simulation, whether the in-loop pressure-volume (P-V) hysteresis loop can serve as a *leading indicator* of fatigue-driven proprioception degradation in the practical shared-manifold topology, where chambers share one supply to cut valve count. We build a transparent, analytically invertible pipeline: a viscoelastic (standard-linear-solid) pneumatic plant producing rate-dependent P-V loops; a canonical fatigue model (Mullins softening with partial recovery, irreversible compliance drift, late leak growth); fatigue-coupled piecewise-constant-curvature (PCC) tip kinematics; a seeded sensor model; and a shared-manifold-versus-isolated-supply network simulator in which inter-chamber cross-talk lives in the actuation-band dynamics. From 2,000 labeled traces (20 actuators $\times$ 5 life stages $\times$ 2 supply topologies $\times$ 2 contact states $\times$ 5 repetitions), split by **actuator identity**, we evaluate static and dynamic pressure-only correctors and three recalibration policies.

We report three findings, including a negative one. **(1)** A predicted advantage of a dynamic (lagged-input) corrector over a static ridge map under the shared manifold **does not materialize**: the cross-talk is real in the network dynamics but second-order, and the dynamic corrector improves pose error by $\approx 0\%$ at physically reasonable parameters. **(2)** The dominant pressure-only proprioception error is instead the fatigue **compliance-scale drift**, which is topology-independent and causes a young-calibrated estimator's curvature error to grow by roughly two orders of magnitude over life. **(3)** The observable P-V loop area is a strong leading indicator of that drift (Pearson $r = 0.885$, 95% CI [0.835, 0.958] on held-out actuators), and a P-V-health-triggered recalibration policy holds pose error within a stated accuracy budget at **60% fewer recalibrations** than an always-on policy (2 vs 5 per actuator), with the trigger threshold selected on training actuators only. **Scope caveat (stated up front):** all 20 actuators are drawn from a single generative model, so "generalization to unseen actuators" tests robustness to that generator, not real device-to-device variation; and absolute pose errors remain sub-millimeter throughout, so the operational value of triggered recalibration grows with tighter accuracy requirements.

1. Introduction

Soft pneumatic grippers are attractive for safe, compliant manipulation, but two monitoring problems limit their deployment. **Fatigue:** silicone actuators fail by cyclic fatigue, and current

practice detects failure only after rupture or via before/after diagnostics [Mosadegh 2014; Wong 2026]. **Proprioception:** recovering pose without added sensors uses the pressure signals already present in the control loop [L. Wang 2020]. Both problems are usually studied in isolation and, more importantly, under an *independent-chamber* assumption.

The most practical multi-finger topology shares a pneumatic manifold across chambers to reduce valve count and onboard mass [McDonald 2021; Moran 2024]. On a shared manifold, driving one chamber's valve perturbs the shared supply node and therefore its neighbors — inter-chamber **cross-talk**. This paper asks whether, in that topology, fatigue and proprioception are *physically coupled*: as a chamber fatigues its compliance rises, which (a) changes the P-V loop shape and (b) could change the cross-talk, so that a single pressure stream might monitor both health and pose.

We test this coupling in a deliberately transparent simulation, before any hardware. The modeling choices are made for analytic checkability (noise-free pose is recoverable to machine precision; isolated supplies give numerically zero cross-talk), and every estimator is validated against its own ground truth. The contribution of this paper is therefore the **pipeline, the design outputs, and a pre-registered modeling result framed as a prediction** — not an experimental claim.

2. Related work and prior-art boundary

P-V hysteresis and fatigue. P-V hysteresis is *already* used as a before/after fatigue diagnostic: Mosadegh et al. [2014] assess PneuNet durability across $>10^6$ cycles by comparing P-V curves, and US Patent 10,639,801B2 [Mosadegh, Shepherd, Whitesides 2020] documents the same before/after protocol. Libby et al. [2023] show the loop shifts with fatigue, and report FEM-agreement degradation (96%→80%). Hysteresis-loop area as a damage proxy has strong cross-domain precedent — metals [Haghshenas 2021], flight-control health indicators [Guo 2021], SHM/acoustic-emission RUL [Galanopoulos 2023] — grounded in early-warning theory [Scheffer 2009] and the prognostics-pipeline / health-indicator-quality criteria of Lei et al. [2018]. **We therefore do not claim first use of P-V for fatigue.** The boundary is the *cycle-resolved, operational leading indicator* with quantified lead behavior, separating irreversible drift from reversible Mullins recovery [Lavazza 2023; Liao 2021; Mars & Fatemi 2002].

Pressure-only proprioception. Multi-chamber pressure-only sensing exists [L. Wang 2020; L. Wang 2023; J. Wang 2025; Joshi & Paik 2023; Zou 2024; Preechayasomboon & Rombokas 2021], but assumes independent or clean per-chamber supply. The nearest collision is Lindenroth et al. [2023], which senses contact force on *coupled parallel* chambers but never frames inter-chamber cross-talk as a proprioception confound or links it to fatigue.

Shared-manifold modeling. Lumped resistance-compliance models of pneumatic circuits [Stanley 2021] provide the tool to model cross-talk as a function of chamber compliance; shared single-source architectures are common in practice [Moran 2024; McDonald 2021].

Correctors and recalibration. ARX/echo-state correctors and physical reservoir computing are established for pneumatic hysteresis [Jaeger & Haas 2004; Youssef 2022; Shen 2025]. Online adaptation to soft-sensor drift exists but is *always-on* [Kushawaha 2025; Thuruthel 2022], or masks anomalous sensors [Sugiyama 2025]; a *fatigue-health-triggered* recalibration is the contrast we evaluate.

Ground truth and evaluation. We follow the soft-proprioception evaluation conventions of Zhang et al. [2023] and report pose RMSE and contact F1; PCC [Webster & Jones 2010] is treated as a modeled output, never as measured truth.

3. Modeling and methods

The pipeline is implemented in Python (sim/, pipeline/, scripts/) with unit-tested gates at each stage. All randomness is seeded; the dataset is reproducible from a fixed seed and its integrity is pinned by a manifest SHA-256.

3.1 Viscoelastic plant and P-V loops (Phase A)

Chamber walls are modeled as a standard linear solid (Zener) in the pressure-volume analog: a spring k_1 in parallel with a Maxwell arm (k_2 , relaxation time τ). This yields rate-dependent P-V hysteresis whose dissipation per cycle has the closed form $E_{\text{loss}}(\omega) = \pi k_2 A^2 (\omega\tau)/(1+(\omega\tau)^2)$, used as a ground-truth check on the numerical loop area, and reduces to the linear cross-talk model in the quasi-static limit. (*Validated: Phase A PASS.*)

3.2 Fatigue model (Phase B)

A canonical, deterministic life law injects: Mullins softening with a 30% permanent floor and 24-h recovery time constant; irreversible compliance drift; curvature-acceleration onset at 70% life; and late leak growth (to 20× conductance) observed by a separate closed-valve pressure-decay probe. Degradation enters the plant as a rising compliance multiplier (softer

k_1) and loss multiplier (larger k_2). (*Validated: Phase B PASS, synthetic consistency.*)

3.3 P-V health features (Phase C)

Causal P-V features, Mullins-recovery normalization, health-index metrics, and matched false-alarm detectors were validated across a 144-condition identifiability grid; a matched segmented model recovers the quadratic onset and fails on logistic-onset variants — the intended inverse-crime/generalization control. (*Validated: Phase C PASS, synthetic.*) Phase C metrics are applied, not re-validated, downstream.

3.4 Fatigue-coupled PCC kinematics (Phase D)

Realized chamber pressure maps to bending curvature through the fatigue compliance multiplier, $\kappa = \kappa_{\text{gain}} \cdot C(\text{life}) \cdot \max(P - P_0, 0)$, so fatigue *amplifies* curvature at fixed pressure. Curvature maps to tip SE(3) by the robot-independent PCC convention [Webster & Jones 2010]. The map is analytically invertible: forward↔inverse round-trips to machine precision, and a noise-free pose is recoverable exactly (gate test).

3.5 Sensor model

A seeded measurement model applies Gaussian noise, quantization, decimation (sampling), and a contact observation (binary flag + synthetic normal force from tip penetration). Independent sub-streams per channel make the corruption reproducible and channel-order-independent.

3.6 Shared-manifold vs isolated network dynamics

A time-domain integrator drives a focal chamber plus two neighbors. In the **shared** topology all chambers draw from one supply node; driving one valve pulls the shared node and perturbs the neighbors. In the **isolated** topology each chamber has its own regulated node and the state blocks are decoupled. Consistent with a lumped-RC pre-test (Gate 0), the cross-talk is *dynamic*: its DC gain is compliance-independent while its actuation-band gain drifts with fatigue. Excitations are multisine in the 1–5 Hz actuation band so the effect is observable. Gate tests confirm: isolated coupling is numerically zero, shared coupling is measurable, and the DC-independent / actuation-band-drift property holds.

3.7 Dataset

20 actuators × 5 life fractions {0.1,0.3,0.5,0.7,0.9} × 2 topologies × 2 contact states × 5 repetitions = **2,000 traces** (240 time samples each). Per-actuator geometry, wall stiffness, and rupture life are drawn deterministically. The split is by **actuator identity** (14 train / 6 held-out), never by trace, so

generalization is measured on actuators never seen in training. Features stored per trace include the shared-manifold pressure, all chamber valve commands, the focal chamber pressure and volume, the true and measured tip positions, the true curvature, and contact.

3.8 Proprioception correctors (Phase E)

The target is focal bending curvature (and, through PCC, tip position) from pressure-only features [shared-manifold pressure, all chamber commands]. We compare a **static ridge** map (memoryless) against a **dynamic** corrector that is the identical ridge estimator on lagged inputs (an exogenous ARX/FIR form; output feedback is excluded because the pose is the estimation target, not an online measurement). The lag count is the only difference.

3.9 P-V health signal and recalibration policies (Phase F)

The health signal is the **observable** P-V loop area from a volumetric probe (the loop-shape quantity, not the ground-truth compliance), normalized to its young value so the trigger is a fractional loop-area growth comparable across actuators. Three policies, each starting from one calibration at the youngest life stage: **fixed** (never recalibrate), **always** (recalibrate every stage), and **triggered** (recalibrate when fractional loop-area growth since the last calibration exceeds τ). The threshold τ is selected on **training actuators only** against a stated accuracy budget, then applied unchanged to held-out actuators. We report estimation error **and** recalibration count together, so a "savings" cannot hide degraded accuracy.

4. Results

4.1 Validation gates

All structural gates pass: noise-free pose recovery and forward $\leftrightarrow$ inverse curvature round-trips to machine precision; isolated-supply off-diagonal coupling at the solver-noise floor ($<10^{-6}$) while shared coupling is measurable ($>10^{-4}$); and the Gate-0 property that DC cross-talk is compliance-independent while actuation-band cross-talk drifts with fatigue. (125 unit tests pass.)

4.2 Cross-talk is real but second-order (negative result, retained; Fig. 1)

The proposal predicted that shared-manifold cross-talk drift would be visible to a dynamic corrector and invisible to a static ridge, making the dynamic corrector win under the shared manifold. **This does not occur.** With all chamber commands available as features and per-actuator calibration, the dynamic (lagged-input) corrector improves curvature RMSE over the static ridge by $\approx 0\%$ **under either topology** (shared: 0.244 vs 0.244 1/m; isolated: 0.295 vs 0.295 1/m; pose ≈ 0.84 vs 0.84 mm and 1.02 vs 1.02 mm). The cross-talk is genuine in the network dynamics (§4.1) but is a small perturbation ($\approx 4\%$ of the focal curvature at these physically reasonable parameters) that does not materially degrade pose estimation. We retain this as a negative result and did not tune network parameters to manufacture the predicted effect; whether a much softer supply enlarges the cross-talk enough to matter is left as an explicit parameter-sensitivity question.




Figure 1. Static vs dynamic (lagged-input) corrector on held-out actuators at old life; the dynamic corrector yields ~0% improvement under either supply topology — the predicted shared-manifold cross-talk advantage does not appear.

4.3 Compliance drift dominates the proprioception error (Fig. 2)

A young-calibrated static estimator's curvature RMSE grows by roughly two orders of magnitude over life ($\approx 0.004 \rightarrow \approx 0.46\text{--}0.56$ 1/m at 0.9 life), and the shared and isolated curves are comparable — the difference does not consistently favor the shared-manifold-degrades-more hypothesis. The dominant pressure-only proprioception error is therefore the **topology-independent fatigue compliance-scale drift**, not cross-talk. This both explains §4.2 and motivates recalibration as the right intervention.




Figure 2. Young-calibrated static estimator curvature RMSE over normalized life (shared vs isolated supply); error grows ~100x by end of life and is comparable across topologies — the degradation is the topology-independent compliance drift, not cross-talk.

4.4 P-V loop area is a leading indicator (Fig. 3)

On held-out actuators, the observable P-V loop-area fractional growth tracks the fixed-calibration pose error over life with **Pearson $r = 0.885$, 95% CI [0.835, 0.958]** (bootstrap, 2,000 resamples; CI excludes 0). The observable loop shape leads the proprioception degradation — the core positive result.




Figure 3. Observable P-V loop-area growth (left axis) leads the fixed-calibration pose error (right axis) over life on held-out actuators; bootstrap $r = 0.885$, 95% CI [0.835, 0.958].

4.5 Recalibration trade-off (Fig. 4)

Against a stated 0.159 mm accuracy budget (selected on training actuators as halfway from the always-on error toward the never-recalibrate error), the trigger threshold $\tau^* = 0.05$ fractional loop-area growth, selected on training actuators, transfers to held-out actuators as:

Policy	Pose RMSE (held-out)	Recalibrations / actuator
fixed (never)	0.40 mm	1
P-V-triggered ($\tau^*=0.05$)	0.06 mm	2
always-on	0.02 mm	5

P-V-triggered recalibration holds pose error near the always-on band (well inside budget) at **60% fewer recalibrations** (2 vs 5), and $\approx 7\times$ better than fixed. Absolute errors are sub-mm in all policies (see §5), so the demonstrated value is the error-vs-cost trade-off and its transfer to unseen actuators, not the absolute accuracy.




Figure 4. Recalibration trade-off on held-out actuators — pose error vs recalibrations per actuator for fixed, P-V-triggered, and always-on policies; the triggered policy meets the accuracy budget at 60% fewer

recalibrations than always-on.

5. Discussion and limitations

- **Simulation only.** Every result is synthetic. The pipeline and the modeling result (framed as a prediction) are the contributions; no experimental validation is claimed.
- **One generative model.** All 20 actuators come from a single generator, so held-out evaluation tests robustness to that generator's variation, not real device-to-device spread. This is the first thing a reviewer should weigh and is stated in the abstract.
- **The headline cross-talk hypothesis was not supported.** Cross-talk is monotone in compliance (Gate 0) but second-order for pose at realistic parameters. The proposal anticipated this "benign coupling" outcome as a publishable floor; the deliverable is the leading-indicator + recalibration result via the compliance-drift pathway.
- **Sub-millimeter absolute errors.** Even never-recalibrating yields ≈ 0.40 mm pose RMSE at this sensor-noise level, so the operational case for triggered recalibration strengthens with tighter pose requirements, softer actuators (larger per-stage drift), or higher noise.
- **Modeling choices.** The fatigue law, PCC single-segment kinematics, and lumped network are transparent assumptions chosen for analytic checkability, not calibrated fits.

6. Conclusion

In simulation, the observable P-V loop shape is a quantified leading indicator ($r \approx 0.89$) of fatigue-driven pressure-only proprioception drift in shared-manifold soft grippers, and a P-V-health-triggered recalibration policy meets an accuracy budget at a fraction of the recalibration cost of always-on adaptation, transferring from training to held-out actuators. The shared-manifold cross-talk, while real, is second-order for pose at realistic parameters — a retained negative result. The natural next step is a physical bench campaign to anchor the synthetic story (single-actuator micro-tear spot-check, then the $N \approx 10$ fatigue campaign of the companion proposal).

7. Reproducibility

All code, the dataset generator, and the analysis scripts are in the repository. The dataset is regenerable from a fixed seed (python -m scripts.phaseD_dataset; integrity pinned by the manifest SHA-256); the studies are python -m scripts.run_study2 (correctors) and python -m scripts.run_study3 (leading indicator + recalibration). Frozen result JSON and figures are committed under data/sim/phaseD/; the pre-registered claim and target figures are in docs/result_spine.md. 125 unit tests gate the pipeline.

References

Citations follow the annotated literature base (docs/A01_A04_Literature_Review.md), where full DOIs are listed. Selected works cited here:

- Mosadegh et al. 2014, *Adv. Funct. Mater.* 24(15):2163–2170. DOI 10.1002/adfm.201303288.
- Mosadegh, Shepherd, Whitesides 2020, US Patent US10639801B2.
- Libby et al. 2023, *ISMR 2023*. DOI 10.1109/ISMR57123.2023.10130227.
- Wong, Luo, Scharff 2026, *Adv. Robotics Research* e202500172. DOI 10.1002/adrr.202500172.
- Mars & Fatemi 2002, *Int. J. Fatigue* 24(9):949–961. DOI 10.1016/S0142-1123(02)00008-7.

- Lavazza, Contino, Marano 2023, *Mech. of Materials* 178:104560. DOI 10.1016/j.mechmat.2023.104560.
- Liao et al. 2021, *Int. J. Mech. Sci.* 206:106624. DOI 10.1016/j.ijmecsci.2021.106624.
- Scheffer et al. 2009, *Nature* 461:53–59. DOI 10.1038/nature08227.
- Lei et al. 2018, *MSSP* 104:799–834. DOI 10.1016/j.ymssp.2017.11.016.
- Haghshenas, Jang, Khonsari 2021, *Mech. of Materials* 155:103734. DOI 10.1016/j.mechmat.2020.103734.
- Guo et al. 2021, *Measurement* 186:110076. DOI 10.1016/j.measurement.2021.110076.
- Galanopoulos et al. 2023, *Eng. Structures* 290:116391. DOI 10.1016/j.engstruct.2023.116391.
- Sugiyama et al. 2025, *Front. Robot. AI* 11:1504651. DOI 10.3389/frobt.2024.1504651.
- L. Wang & Z. Wang 2020, *Soft Robotics* 7(2):198–217. DOI 10.1089/soro.2018.0135.
- L. Wang et al. 2023, *Soft Robotics* 10(4):825–837. DOI 10.1089/soro.2021.0056.
- J. Wang et al. 2025, *Adv. Intell. Syst.* 7(4):2400534. DOI 10.1002/aisy.202400534.
- Joshi & Paik 2023, *Soft Matter* 19. DOI 10.1039/D2SM01197B.
- Zou et al. 2024, *Nature Communications* 15:539. DOI 10.1038/s41467-023-44517-z.
- Preechayasomboon & Rombokas 2021, *Actuators* 10(2):30. DOI 10.3390/act10020030.
- Lindenroth et al. 2023, *IEEE/ASME T. Mechatronics* 28(1):80–91. DOI 10.1109/TMECH.2022.3210065.
- Moran et al. 2024, *Comms. Eng.* 3:117. DOI 10.1038/s44172-024-00251-y.
- McDonald & Ranzani 2021, *Front. Robot. AI* 8:720702. DOI 10.3389/frobt.2021.720702.
- Stanley et al. 2021, *ASME J. Dyn. Sys. Meas. Control* 143(5):051001. DOI 10.1115/1.4049009.
- Jaeger & Haas 2004, *Science* 304(5667):78–80. DOI 10.1126/science.1091277.
- Youssef et al. 2022, *Micromachines* 13(2):216. DOI 10.3390/mi13020216.
- Shen, Miyazaki, Kawashima 2025, *RA-L*, arXiv:2409.06961.
- Kushawaha et al. 2025, arXiv:2503.16540.
- Thuruthel, Gardner, Iida 2022, *Soft Robotics* 9(6):1167–1176. DOI 10.1089/soro.2021.0012.
- Webster & Jones 2010, *IJRR* 29(13):1661–1683. DOI 10.1177/0278364910368147.
- Zhang et al. 2023, *IROS 2023*:2564–2571. DOI 10.1109/IROS55552.2023.10342379.